

# RandomGen - Weighted Random Number Generator

---



## Project Overview

A production-ready implementation of a random number generator with specified probabilities. It allows you to define a set of possible values and their respective probabilities of occurrence, then efficiently generate random numbers that follow this distribution.

## Features

- **High Performance:**  $O(\log n)$  lookup algorithm using binary search on cumulative probabilities
- **Thread-Safe:** Safe for concurrent use from multiple threads
- **Well-Tested:** Comprehensive test suite with statistical validation
- **Production-Ready:** Robust error handling, documentation, and CI/CD pipeline
- **Cross-Platform:** Works on Windows, macOS, and Linux

## Implementation Details

### Core Algorithm

The RandomGen class uses the following approach to generate weighted random numbers:

#### 1. Initialization:

- Validate input arrays (sizes match, probabilities sum to 1, etc.)
- Precompute cumulative probabilities using `std::partial_sum`
- Set up thread-local random number generator

#### 2. Random Number Generation:

- Generate a uniform random value between 0 and 1
- Use binary search (`std::lower_bound`) to find the appropriate index in the cumulative probability array
- Return the corresponding value from the randomNums array

### Time Complexity

- **Construction:**  $O(n)$  where  $n$  is the number of possible values
- **nextNum():**  $O(\log n)$  due to binary search on the cumulative probabilities
- **Memory Usage:**  $O(n)$  to store the input arrays and cumulative probabilities

### Thread Safety

RandomGen is thread-safe by design:

- Uses `thread_local` for the random number generator to avoid contention
- All member variables are initialized once and become immutable after construction

- No internal state is modified during calls to `nextNum()`

## Getting Started

### Option 1: Using Precompiled Binaries

Precompiled binaries for Windows, macOS, and Linux are available on the [Releases](#) page.

1. Download the appropriate zip file for your platform
2. Extract the contents
3. Run the example or test executable:

```
# Run example
./RandomGenExample

# Run tests
./RandomGenTests
```

### Option 2: Using Docker

If you have Docker installed, you can run RandomGen without any other dependencies:

```
# Pull the image
docker pull ghcr.io/ayanchev01/randomgen:latest

# Run the example
docker run --rm ghcr.io/ayanchev01/randomgen:latest example

# Run the tests
docker run --rm ghcr.io/ayanchev01/randomgen:latest tests

# Run both
docker run --rm ghcr.io/ayanchev01/randomgen:latest all
```

### Option 3: Building From Source

#### Prerequisites

- C++20 compatible compiler (GCC 10+, Clang 10+, MSVC 2019+)
- CMake 3.15 or higher
- Git (optional, for retrieving GoogleTest)

#### Build Instructions

Clone the repository:

```
git clone https://github.com/AYanchev01/RandomGen.git
cd RandomGen
```

Build the project:

```
mkdir build
cd build
cmake .. -DCMAKE_BUILD_TYPE=Release
cmake --build . --config Release
```

Run the example:

```
./bin/Release/RandomGenExample
```

Run the tests:

```
./bin/Release/RandomGenTests
```

## API Documentation

### RandomGen Class

```
class RandomGen {
public:
    /**
     * @brief Construct a new RandomGen object
     *
     * @param randomNums Values that may be returned by nextNum()
     * @param probabilities Probability of occurrence for each value
     * @throw std::invalid_argument if inputs are invalid
     */
    RandomGen(const std::vector<int>& randomNums, const std::vector<double>&
probabilities);

    /**
     * @brief Returns one of the randomNums based on their probabilities
     *
     * When called multiple times over a long period, it returns numbers
     * roughly with the initialized probabilities.
     *
     * This method is thread-safe and can be called concurrently from multiple
     threads.
     *
     * @return int A randomly selected number
     */
};
```

```
    */  
    int nextNum() noexcept;  
};
```

## Usage Example

```
#include "RandomGen.h"  
#include <iostream>  
#include <vector>  
  
int main() {  
    // Define possible values and their probabilities  
    std::vector<int> randomNums = {-1, 0, 1, 2, 3};  
    std::vector<double> probabilities = {0.01, 0.3, 0.58, 0.1, 0.01};  
  
    // Create random generator  
    RandomGen randomGen(randomNums, probabilities);  
  
    // Generate and print random numbers  
    for (int i = 0; i < 10; ++i) {  
        std::cout << randomGen.nextNum() << " ";  
    }  
    std::cout << std::endl;  
  
    return 0;  
}
```

## Testing

The project includes a comprehensive test suite that validates:

### 1. Input Validation:

- Handling of invalid probability distributions (not summing to 1)
- Handling of negative probabilities
- Handling of mismatched array lengths
- Handling of empty arrays
- Handling of edge cases (probabilities very close to 1)

### 2. Functionality:

- Single value with probability 1
- Uniform distribution behavior
- Distribution matching specified probabilities
- Statistical validation using confidence intervals

### 3. Thread Safety:

- Concurrent access from multiple threads

# Continuous Integration/Continuous Delivery

The project includes a full CI/CD pipeline implemented with GitHub Actions:

## Continuous Integration

- **Build Matrix:** Builds and tests on Windows, macOS, and Linux with both Debug and Release configurations
- **Code Quality:** Static analysis with cppcheck and clang-tidy
- **Code Coverage:** Measures test coverage and uploads to Codecov

## Continuous Delivery

- **Cross-platform Builds:** Automated builds for Windows, macOS (universal binary for Intel and Apple Silicon), and Linux
- **Release Automation:** Creates GitHub releases with precompiled binaries
- **Docker Images:** Builds and pushes Docker images to GitHub Container Registry